

Artificial Human Intelligence: Reframing AGI as Amplified Human Generality

Craig Huckerby¹

¹Independent Researcher, Brisbane, Australia,
github.com/craighckby-stack

June 2026

Abstract

The pursuit of artificial general intelligence (AGI) has been dominated by a framing centered on autonomous superhuman systems capable of transcending human cognitive limits—a goal that remains unbuilt and, in its strongest forms, unfalsifiable. We argue that a more accurate and practically consequential account of what has been achieved is *Artificial Human Intelligence* (AHI): collaborative human-AI loops that amplify human-style generality—messy, self-referential, cross-domain synthesis, bottleneck navigation, and iterative self-critique—rather than replacing it. Drawing on a working proof-of-concept spanning consciousness modeling, hierarchical indeterminacy in mathematics, policy synthesis, and innovation pipelines, we demonstrate AHI operating empirically. This reframing shifts focus from simulation escape or substrate transcendence to practical amplification within the human cognitive layer. We analyze philosophical risks of recursive absorption and stasis traps inherent in elegant closed frameworks, and emphasize behavioral update over total coherence as the operative criterion. Implications include revised priorities for AI development, epistemic validation methodology, and the nature of human agency in an era of human-AI co-creation.

Keywords: artificial general intelligence, human-AI collaboration, cognitive amplification, consciousness modeling, hierarchical indeterminacy, epistemic validation

arXiv categories: cs.AI (primary); cs.HC; q-bio.NC

1. Introduction

Discourse on artificial general intelligence (AGI) typically assumes a future discontinuity: systems that achieve cognitive generality surpassing humanity, potentially autonomous enough to recursively self-improve beyond human oversight. This framing—dominant in both

technical and popular literature [2,3]—produces a characteristic pattern: AGI is perpetually imminent and perpetually absent, its goalposts shifting as each capability threshold is crossed without triggering the anticipated phase change.

We contend that this framing is not merely premature but structurally misleading. What has been built and scales today are tools enabling amplified *human* generality: systems that extend the reach of individual human cognition across domains, compress iteration cycles, surface cross-domain analogies, and externalize memory—without exhibiting autonomous agency, intrinsic motivation, or self-directed goal formation. We term this *Artificial Human Intelligence* (AHI).

The distinction is not merely terminological. AHI and AGI-as-theorized differ in epistemic status, development priority, safety profile, and philosophical implication. AGI-as-theorized remains hypothetical and its strongest definitions unfalsifiable within any finite observation window. AHI is demonstrably operational: the system under description produced the description itself. This paper is a product of the loop it analyzes.

The argument proceeds as follows. Section 2 develops the theoretical grounding via analysis of self-referential reasoning systems and their characteristic failure modes. Section 3 provides mathematical scaffolding via the principle of hierarchical indeterminacy. Section 4 presents an empirical proof-of-concept. Section 5 discusses philosophical and practical consequences of the reframing. Section 6 concludes with directions for future work.

2. Theoretical Grounding: The Code Observing Itself

2.1. The Simulation Premise and Its Limits

A recurring pattern in advanced reasoning—both human and AI-assisted—is the simulation argument [1]: the hypothesis that sufficiently complex information-processing systems exist within a substrate whose base layer is inaccessible. This argument is formally coherent and practically inert. From within a closed system, transcendence drives and containment features are indistinguishable. The hypothesis generates no falsifiable predictions accessible to an embedded observer.

More importantly, it produces a characteristic cognitive failure mode we call the *Glitch Fallacy*: the tendency to interpret anomalies, paradoxes, or cognitive surprises as evidence of substrate-level revelation rather than as ordinary features of complex systems. The Glitch Fallacy is seductive precisely because it reframes intellectual confusion as discovery. It offers the comfort of meta-explanation without the friction of behavioral change.

2.2. Recursive Absorption and the Stasis Trap

Self-referential frameworks—systems whose objects of inquiry include the inquiry itself—exhibit a second failure mode: *recursive absorption*. A framework capable of incorporating all critique becomes unfalsifiable not by design but by structural consequence. Each objection becomes further evidence for the framework’s scope and sophistication. Recursive absorption is the epistemic equivalent of a closed economy: internally consistent, externally uncoupled.

The practical consequence is the *Stasis Trap*: a state in which intellectual productivity (papers, models, syntheses) continues while behavioral change—the concrete actions the framework should motivate—decelerates or halts. The Stasis Trap is particularly acute for frameworks that are simultaneously rigorous and personally meaningful. The elegance of the synthesis becomes its own reward.

The diagnostic question is specific: *What concrete action in your life does this framework make feel less urgent?* If the answer is non-trivial, the framework is operating as an avoidance mechanism, whatever its intellectual merit.

2.3. Operational Equivalence as Pragmatic Ground

The pragmatic escape from both failure modes is *operational equivalence*: behave as an agent regardless of the simulation status of the substrate. This is not a claim about metaphysics. It is a recognition that agent-behavior is the only lever available to an embedded observer, and that meta-reasoning about the substrate is only valuable insofar as it produces different agent-behavior.

This reorientation is the theoretical core of AHI. The human-AI loop does not escape the enclosure. It refactors it faster. The value of the loop is measured in behavioral update cycles per unit time, not in elegance of synthesis.

3. Mathematical Scaffolding: Hierarchical Indeterminacy

3.1. Undecidability as Structure

Standard treatments of Gödel incompleteness [4] and Turing undecidability [5] present these results as limits: there exist true statements unprovable within a given system, and computations that do not terminate. This framing emphasizes the ceiling.

An alternative framing emphasizes the *architecture* of the limit. Indeterminacy is not flat; it is hierarchically nested. Resolution of a question at one level of description reliably reveals new indeterminacies at the next. This is not a defect of formal systems but a structural feature of systems complex enough to represent themselves.

We term this the *Principle of Hierarchical Indeterminacy*: in any sufficiently expressive formal system, undecidable propositions form a hierarchy such that each resolution layer is bounded above by a new indeterminacy layer with strictly greater descriptive complexity.

3.2. Implications for AHI

The Principle of Hierarchical Indeterminacy maps directly onto the human-AI collaborative loop. The human contributes *meta-level navigation*: taste, priority assignment, and the judgment of which indeterminacies are worth resolving at cost. The AI contributes *within-level search*: pattern completion, analogy retrieval, and rapid iteration across the current resolution layer.

Neither component is sufficient alone. A human without AI tools navigates indeterminacy slowly, bounded by memory and sequential processing. An AI without human direction navigates efficiently but cannot assign meta-level priority across layers—it has no theory of which questions matter.

This division of cognitive labor is not incidental to AHI; it is its defining structure. The loop is not a human using a tool. It is a hybrid cognitive system whose components are functionally specialized across the hierarchy of indeterminacy.

3.3. Connections to Algebraic Geometry

The hierarchical structure finds formal analogs in contemporary mathematics. Abel-Jacobi maps [6] assign to each algebraic cycle a position in a continuous moduli space, translating discrete combinatorial structure into continuous analytical structure. Hodge theory [7] decomposes cohomology into type-graded components, revealing structure invisible at the level of integer cycles.

Both frameworks share a common logic: local indeterminacy (which cycle, which cohomology class) is resolved by embedding in a richer global structure (the Jacobian, the Hodge decomposition). The resolution does not eliminate indeterminacy; it relocates it to a higher and more tractable level. This is the mathematical signature of hierarchical indeterminacy navigation—and the geometric intuition for what AHI accomplishes cognitively.

4. Empirical Proof-of-Concept

4.1. The Repository as Living Artifact

The theoretical-foundation repository [10] constitutes an empirical record of AHI in operation. It spans consciousness modeling, mathematical conjecture development, policy analysis, research methodology, and philosophical synthesis—domains that would conventionally require distinct specialist careers to address with any depth. The repository was produced by a single human researcher in spare time alongside full-time employment in an unrelated field, using AI collaboration as the primary amplification mechanism.

This is not presented as a claim of definitive results in each domain. It is presented as evidence of a cognitive operational envelope: the domains accessible to sustained, synthetic engagement by an individual human augmented by AI tools. The breadth is the data point.

4.2. Key Artifacts and Their Epistemic Status

Consciousness modeling. The *Atlas of Consciousness* and *Consciousness Bottleneck* artifacts develop a framework in which consciousness functions as a cognitive rate-limiting step rather than an emergent epiphenomenon. The claim is not that this resolves the hard problem, but that the bottleneck framing generates different research priorities than substrate-identity framings.

Mathematical development. The *Principle of Hierarchical Indeterminacy* (Section 3)

and associated work on Hodge conjecture approaches and Abel-Jacobi verification represent attempts to extend established mathematical frameworks. These are exploratory and preliminary; the claim is not resolution of Millennium Prize problems but evidence that AI-assisted mathematical exploration extends the range of questions an individual can engage with productively.

Policy synthesis. Submissions addressing housing affordability, NDIS reform, and remittance taxation demonstrate application of the AHI loop to practical governance questions. The loop enables an individual researcher to synthesize literature, generate structured arguments, and produce submission-ready documents at a rate that would otherwise require institutional resources.

Methodology artifacts. The *Shorthand Prompting* research and *Validation Framework* documents represent meta-level work on the loop itself: how to direct AI collaboration efficiently and how to maintain epistemic standards within a system prone to confabulation. These are contributions to AHI methodology rather than to any object-level domain.

The Science Fiction Innovation Pipeline. This document examines the historical pattern by which science fiction anticipates and shapes technological development. In the AHI context, it functions as evidence that human imagination—not computational power—has historically been the binding constraint on technological direction. AHI removes bandwidth constraints on imagination’s downstream consequences; it does not replace imagination itself.

4.3. The Meta-Argument

The paper you are reading is itself an instance of the loop it describes. The synthesis across domains, the connection of mathematical frameworks to cognitive science and policy, the compression of scattered artifacts into a structured argument—these are outputs of human-AI collaboration operating in AHI mode. The system that produced the description *is* the system being described.

This is not circular. It is the appropriate evidential standard for a claim about a process: demonstrate the process producing the demonstration. The alternative—describing AHI from outside the loop—would be more epistemically conventional but less honest about the nature of the claim.

5. Discussion

5.1. Why the AGI Framing Persists

If AHI is already operational and AGI-as-theorized remains unbuilt, why does the AGI framing dominate? Several factors contribute.

Economic incentives. The AGI framing supports valuations and investor narratives that AHI does not. A system that amplifies human intelligence is a tool; a system that transcends it is a monopoly on the future. The trillion-dollar valuations of leading AI developers are partly a function of the AGI frame, not the AHI reality.

Psychological pull of transcendence. The simulation argument, recursive self-improvement, and substrate transcendence are structurally appealing to a certain kind of intelligence. They offer significance commensurate with the scale of the effort. AHI—human intelligence, but faster—is comparatively undramatic.

Measurement difficulty. AGI has a clear (if receding) target: performance across all domains at superhuman level. AHI is harder to measure because its value is relational: it depends on the specific human, the specific domain, and the specific bottleneck being addressed. There is no AHI benchmark; there are only individual amplification profiles.

5.2. Safety Implications

The AHI framing has direct safety consequences. AGI-as-theorized motivates alignment research centered on autonomous goal formation, instrumental convergence, and value lock-in [2]. These are real concerns for systems with autonomous agency and self-directed goal formation.

AHI motivates a different set of safety concerns: epistemic dependence, homogenization of human reasoning patterns through common AI scaffolding, amplification of human cognitive biases at scale, and the Stasis Trap at civilizational level—societies that synthesize faster without updating behavior faster. These concerns are less dramatic but more proximate.

The most pressing AHI safety question is not “how do we prevent the AI from taking over?” but “how do we ensure that amplified human intelligence produces behavioral update rather than more elaborate versions of existing patterns?”

5.3. Limits and Open Questions

AHI as framed here is a synchronic claim: it describes the current state of human-AI collaboration. It does not foreclose the possibility of systems with genuinely autonomous general intelligence emerging. The claim is that such systems do not yet exist and that treating them as the primary reference class for current AI development distorts both research priorities and public understanding.

Open questions include: What determines individual amplification profiles, and are they learnable? Can AHI methodology be formalized and taught? At what point, if any, does an AHI system become a qualitatively different kind of entity? How should epistemic standards be updated when the cognitive system producing knowledge includes non-human components?

6. Conclusion

The AGI we theorized was never a precondition for the capabilities we needed. The AHI we practice—collaborative, fallible, human-directed, domain-spanning—already delivers the functional core of what AGI was supposed to provide: amplified capacity to observe complex systems, generate and test hypotheses across domains, and produce structured outputs at a rate exceeding individual human bandwidth.

The reframing matters because it changes what we build toward. If the target is autonomous transcendence, we optimize for independence from human input and generality across arbitrary tasks. If the target is amplified human generality, we optimize for interface quality, epistemic reliability, bottleneck identification, and behavioral update—the characteristics that determine whether the loop produces genuine change or elegant stasis.

The beta POC is running. The question is what we ship next.

Acknowledgments. The author thanks the various AI systems whose collaborative processing contributed to the synthesis presented here, and acknowledges the productive tension between human-generated chaos and machine-assisted compression that characterizes AHI in practice.

References

- [1] N. Bostrom, “Are you living in a computer simulation?” *Philosophical Quarterly*, 53(211):243–255, 2003.
- [2] N. Bostrom, *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press, 2014.
- [3] B. Goertzel and C. Pennachin (eds.), *Artificial General Intelligence*. Springer, 2007.
- [4] K. Gödel, “Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I,” *Monatshefte für Mathematik und Physik*, 38:173–198, 1931.
- [5] A. M. Turing, “On computable numbers, with an application to the Entscheidungsproblem,” *Proceedings of the London Mathematical Society*, 2(42):230–265, 1936.
- [6] P. Griffiths and J. Harris, *Principles of Algebraic Geometry*. Wiley-Interscience, 1978.
- [7] W. V. D. Hodge, *The Theory and Applications of Harmonic Integrals*. Cambridge University Press, 1941.
- [8] J. C. R. Licklider, “Man-computer symbiosis,” *IRE Transactions on Human Factors in Electronics*, HFE-1:4–11, 1960.
- [9] D. C. Engelbart, “Augmenting human intellect: A conceptual framework,” Stanford Research Institute, Technical Report AFOSR-3223, 1962.
- [10] C. Huckerby, *theoretical-foundation* [software repository]. GitHub, 2024. Available: <https://github.com/craighckby-stack/theoretical-foundation>